



# Beam SDK Tutorial

Apache Beam Summit - Beam Introduction

Berlin, June, 2019



# 01 Introduction


The Apache Beam programming model

# What is part of Apache Beam?

## One Model



## Multiple Modes



Batch



Streaming


## Multiple SDKs



Java



Python



Go

## Multiple Runners



**Direct:** local  
for testing



**Cloud Dataflow:** fully  
managed service on  
Google Cloud




**Apache Flink:** local,  
on-premise, cloud



**Apache Spark:** local,  
on-premise, cloud

# What is a pipeline?



- A Directed Acyclic Graph of data **transformations** applied to one or more **collections** of data
- Possibly **unbounded collections** of data flow on the edges
- May include multiple sources and multiple *sinks*
- Optimized and executed as a unit


# Apache Beam Ecosystem



## Beam

- Unified programming model
- Portable
  - Multiple Runners
  - Multiple Languages (Java, Python, Go)
- Extensible:
  - IO: GCP + major open source + APIs
  - DSLs: SQL, Scala

# What is a pipeline?



- Beam represents datasets using an abstraction called PCollection
- Data transformations are represented by an abstraction called PTransform

# The pipeline describes...

**What** are you computing?

**Where** in event time?

**When** in processing time?

**How** do refinements relate?

The pipeline describes...

What = Transformations

Where = Windowing

When = Watermarks + Triggers

How = Accumulation







## 02 Writing a pipeline

**What** results are calculated?


# What are you computing?



**Element-Wise  
(map)**




**Aggregating  
(reduce)**




**Composite  
(reusable combinations)**


# What are you computing?



**Element-Wise  
(map)**



**Aggregating  
(reduce)**




**Composite  
(reusable combinations)**

# Element-wise transforms: ParDo

(ParDo = “Parallel Do”)

Performs a user-provided transformation on each element of a PCollection independently

ParDo can be used for many different operations...



# Element-wise transforms: ParDo (Java)

```
Pipeline p = Pipeline.create(options);  
  
PCollection<String> input = p.apply(...);  
  
firstLetters = input.apply(ParDo.of(  
    new DoFn<String, KV<Char, String>>() {  
        @ProcessElement  
        public void processElement(  
            @Element String word, OutputReceiver<> o) {  
                Char firstLetter = word.charAt(0);  
                o.output(KV.of(word.charAt(0), word));  
            }  
        }  
    ));
```

{Storm, Flink, Apex, Spark, ...}



**ParDo(KeyByFirstLetter)**



{KV<S, Storm>, KV<F, Flink>,  
KV<A, Apex>, KV<S, Spark>, ...}

# Element-wise transforms: ParDo (Python)

```
class FirstLetter(beam.DoFn):  
    def process(self, word):  
        return [word[0]]  
input = ...;  
firstLetters = input | beam.ParDo(FirstLetter)
```

{Storm, Flink, Apex, Spark, ...}



**ParDo(KeyByFirstLetter)**



{KV<S, Storm>, KV<F, Flink>,  
KV<A, Apex>, KV<S, Spark>, ...}

# Element-wise transforms: ParDo (Go)

```
func firstLetter(w: string) string {  
    return w[0]  
}  
lines := ...  
firstLetters := beam.ParDo(s, firstLetter, line)
```

{Storm, Flink, Apex, Spark, ...}




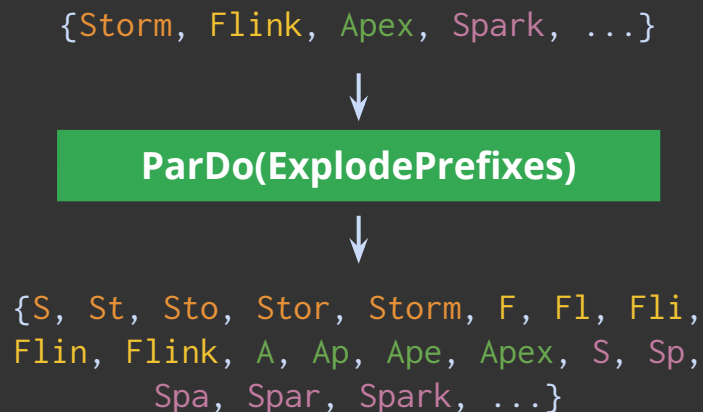
**ParDo(KeyByFirstLetter)**



{KV<S, Storm>, KV<F, Flink>,  
KV<A, Apex>, KV<S, Spark>, ...}

# Element-wise transforms: ParDo

ParDo can output 1, 0 or many values for each input element





# Element-wise transforms: Friends of ParDo

The SDK includes other Element Wise Transforms for convenience

<b>ParDo</b>	General; 1-input to (0,1,many)-outputs; side-inputs and side-outputs
<b>Filter</b>	1-input to (0 or 1)-outputs
<b>MapElements</b>	1-input to 1-output
<b>FlatMapElements</b>	1-input to (0,1,many)-output
<b>WithKeys</b>	value -> KV(f(value), value)
<b>Keys</b>	KV(key, value) -> key
<b>Values</b>	KV(key, value) -> value

# Element-wise transforms: Friends of ParDo

The SDK includes other Element Wise Transforms for convenience

<b>ParDo</b>	General; 1-input to (0,1,many)-outputs; side-inputs and side-outputs
<b>Filter</b>	1-input to (0 or 1)-outputs
<b>MapElements</b>	1-input to 1-output
<b>FlatMapElements</b>	1-input to (0,1,many)-output
<b>WithKeys</b>	value -> KV(f(value), value)
<b>Keys</b>	KV(key, value) -> key
<b>Values</b>	KV(key, value) -> value

```
// Example filter Java
input.apply(Filter
    .byPredicate((String w) -> w.startsWith("S")));
```

# Element-wise transforms: Friends of ParDo

The SDK includes other Element Wise Transforms for convenience

<b>ParDo</b>	General; 1-input to (0,1,many)-outputs; side-inputs and side-outputs
<b>Filter</b>	1-input to (0 or 1)-outputs
<b>MapElements</b>	1-input to 1-output
<b>FlatMapElements</b>	1-input to (0,1,many)-output
<b>WithKeys</b>	value -> KV(f(value), value)
<b>Keys</b>	KV(key, value) -> key
<b>Values</b>	KV(key, value) -> value

```
# Example filter, Python
input | beam.Filter(lambda w: w[0] == 'S')
```

# Element-wise transforms: Friends of ParDo

The SDK includes other Element Wise Transforms for convenience

<b>ParDo</b>	General; 1-input to (0,1,many)-outputs; side-inputs and side-outputs
<b>Filter</b>	1-input to (0 or 1)-outputs
<b>MapElements</b>	1-input to 1-output
<b>FlatMapElements</b>	1-input to (0,1,many)-output
<b>WithKeys</b>	value -> KV(f(value), value)
<b>Keys</b>	KV(key, value) -> key
<b>Values</b>	KV(key, value) -> value

```
// Example filter, Go
filter.Include(s, input, func(s string) bool {
    return w[0] == 'S'
})
```

# Element-wise transforms: Friends of ParDo

The SDK includes other Element Wise Transforms for convenience

<b>ParDo</b>	General; 1-input to (0,1,many)-outputs; side-inputs and side-outputs
<b>Filter</b>	1-input to (0 or 1)-outputs
<b>MapElements</b>	1-input to 1-output
<b>FlatMapElements</b>	1-input to (0,1,many)-output
<b>WithKeys</b>	value -> KV(f(value), value)
<b>Keys</b>	KV(key, value) -> key
<b>Values</b>	KV(key, value) -> value

```
// MapElements Java
input.apply(MapElements
    .into(TypeDescriptors.kvs(TypeDescriptors.characters(),
                             TypeDescriptors.strings()))
    .via((String w) -> KV.of(w, w.charAt(0)))
```

# Element-wise transforms: Friends of ParDo

The SDK includes other Element Wise Transforms for convenience

<b>ParDo</b>	General; 1-input to (0,1,many)-outputs; side-inputs and side-outputs
<b>Filter</b>	1-input to (0 or 1)-outputs
<b>MapElements</b>	1-input to 1-output
<b>FlatMapElements</b>	1-input to (0,1,many)-output
<b>WithKeys</b>	value -> KV(f(value), value)
<b>Keys</b>	KV(key, value) -> key
<b>Values</b>	KV(key, value) -> value

```
# MapElement Python
input | beam.Map(lambda w: (w, w[0]))
```

# Element-wise transforms: Friends of ParDo

The SDK includes other Element Wise Transforms for convenience

<b>ParDo</b>	General; 1-input to (0,1,many)-outputs; side-inputs and side-outputs
<b>Filter</b>	1-input to (0 or 1)-outputs
<b>MapElements</b>	1-input to 1-output
<b>FlatMapElements</b>	1-input to (0,1,many)-output
<b>WithKeys</b>	value -> KV(f(value), value)
<b>Keys</b>	KV(key, value) -> key
<b>Values</b>	KV(key, value) -> value

```
// FlatMapElements Java
input.apply(FlatMapElements
    .into(TypeDescriptors.strings())
    .via((String w) -> populateSuffixes(w)));
```

# Element-wise transforms: Friends of ParDo

The SDK includes other Element Wise Transforms for convenience

<b>ParDo</b>	General; 1-input to (0,1,many)-outputs; side-inputs and side-outputs
<b>Filter</b>	1-input to (0 or 1)-outputs
<b>MapElements</b>	1-input to 1-output
<b>FlatMapElements</b>	1-input to (0,1,many)-output
<b>WithKeys</b>	value -> KV(f(value), value)
<b>Keys</b>	KV(key, value) -> key
<b>Values</b>	KV(key, value) -> value

```
# FlatMapElement Python
input | beam.FlatMap(populateSuffixes)
```



# Element-wise transforms: Friends of ParDo

The SDK includes other Element Wise Transforms for convenience

<b>ParDo</b>	General; 1-input to (0,1,many)-outputs; side-inputs and side-outputs
<b>Filter</b>	1-input to (0 or 1)-outputs
<b>MapElements</b>	1-input to 1-output
<b>FlatMapElements</b>	1-input to (0,1,many)-output
<b>WithKeys</b>	value -> KV(f(value), value)
<b>Keys</b>	KV(key, value) -> key
<b>Values</b>	KV(key, value) -> value

```
// WithKeys Java
input.apply(WithKeys.
  .of((String w) -> w.charAt(0))
  .withKeyType(TypeDescriptors.characters()))
```

# Element-wise transforms: Friends of ParDo


The SDK includes other Element Wise Transforms for convenience

<b>ParDo</b>	General; 1-input to (0,1,many)-outputs; side-inputs and side-outputs
<b>Filter</b>	1-input to (0 or 1)-outputs
<b>MapElements</b>	1-input to 1-output
<b>FlatMapElements</b>	1-input to (0,1,many)-output
<b>WithKeys</b>	value -> KV(f(value), value)
<b>Keys</b>	KV(key, value) -> key
<b>Values</b>	KV(key, value) -> value


```
// Keys
input.apply(Keys.create())
```

```
// Values
input.apply(Values.create())
```


# What are you computing?



**Element-Wise  
(map)**




**Aggregating  
(reduce)**



**Composite  
(reusable combinations)**

# Grouping transforms: GroupByKey

Takes a PCollection of key-value pairs and groups all values with the same key



How can we use GroupByKey to compute the most common value for each key?

# Grouping transforms: GroupByKey

Takes a PCollection of key-value pairs and groups all values with the same key

```
{KV<S, Storm>, KV<F, Flink>, KV<A, Apex>, KV<S, Spark>, ...}
```

**GroupByKey**

```
input.apply(GroupByKey.<Character, String>create())
```

```
{..., ...]>, KV<A, [Apex, ...]>, ...}
```

How can we get the value for each key?

# Grouping transforms: GroupByKey

Takes a PCollection of key-value pairs and groups all values with the same key

```
{KV<S, Storm>, KV<F, Flink>, KV<A, Apex>, KV<S, Spark>, ...}
```

**GroupByKey**

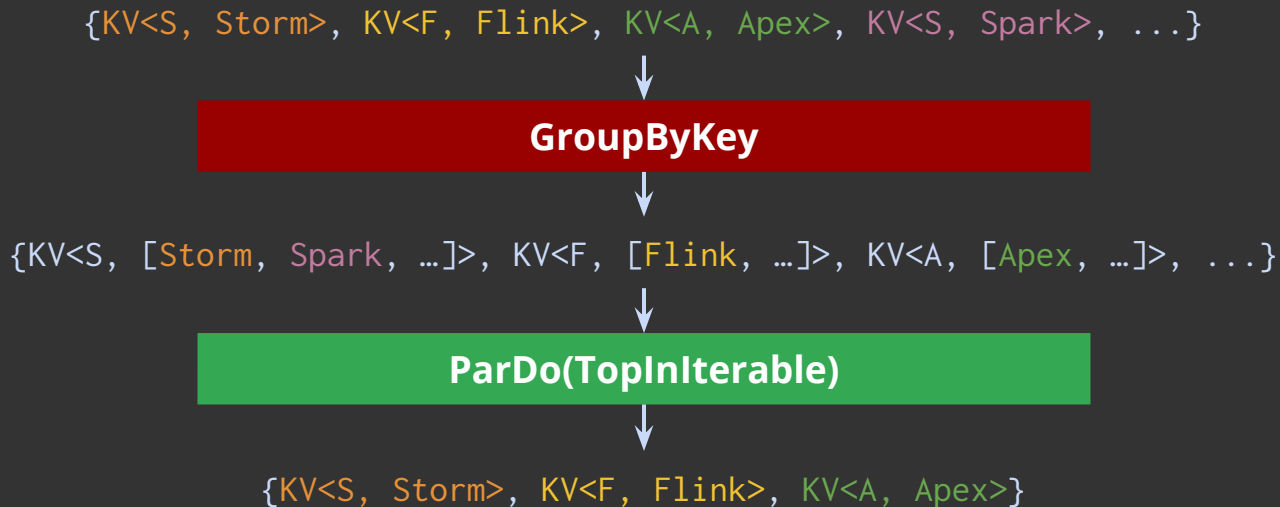
```
input | beam.GroupByKey()
```

```
{..., ...}>, KV<A, [Apex, ...]>, ...}
```

How can we group values by key? on value for each key?

# Grouping transforms: GroupByKey

Computing the most common value for each key



`TopInIterable` processes `KV<K, Iterable<String>>` and has to look at all of the values for each key...

# Grouping transforms: GroupByKey

GroupByKey followed by ParDo can often be simplified (and optimized!): Combine

```
{KV<S, Storm>, KV<F, Flink>, KV<A, Apex>, KV<S, Spark>, ...}
```



**Combine.perKey(CountAndCompare)**



```
{KV<S, Storm>, KV<F, Flink>, KV<A, Apex>}
```



# Grouping transforms: Combine

CountAndCompare is a CombineFn that counts words and then extracts the top-K. You can write your own for any operation that is associative & commutative.



Initialize accumulators




Add Input to each accumulator



Merge accumulators



Merge accumulators (again)



Extract output (from accumulator)

# Grouping transforms: Built-in CombineFns

The SDK includes many pre-defined Combiners:

**Top.perKey(1)**

**Min.longsPerKey()**

**Count.perKey()**

**Max.longsPerKey()**


**Sum.longsPerKey()**

**Mean.longsPerKey()**


**ApproximateQuantiles.perKey(5)**

**ApproximateUnique.perKey(10)**


# Writing a pipeline = Gluing pieces together



**Element-Wise  
(map)**



**Aggregating  
(reduce)**



**Composite  
(reusable combinations)**

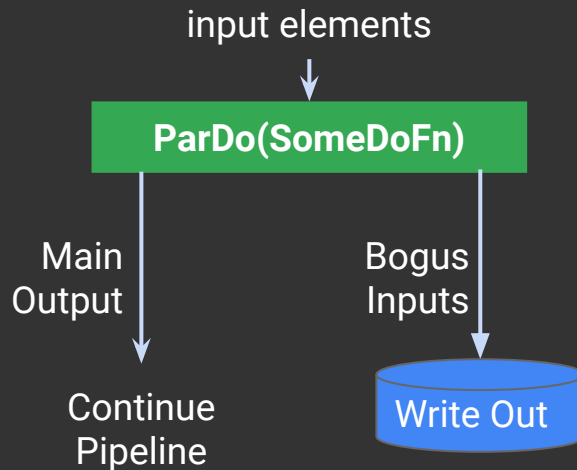
# Multiple outputs

ParDos can produce multiple outputs

Example usage: dead-letter pattern

A main output containing all the successfully processed results

A secondary output containing all the elements that failed to be processed



# Example: Multiple outputs (Java)

```
final TupleTag<Output> successTag = new TupleTag<>() {};  
final TupleTag<Input> deadLetterTag = new TupleTag<>() {};
```

```
PCollection<Input> input = ...;  
PCollectionTuple outputTuple = input.apply(ParDo  
    .withOutputTags(successTag, TupleTagList.of(deadLetterTag))  
    .of(new DoFn<Input, Output>() {  
        @ProcessElement  
        public void processElement(@Element InputT e, MultiOutputReceiver o) {  
            try {  
                o.output(successTag, validateElement(e));  
            } catch (Exception e) {  
                o.output(deadLetterTag, e);  
            }  
        }  
    }));  
PCollection<Output> success = outputTuple.get(successTag);  
PCollection<Input> deadLetters = outputTuple.get(deadLetterTag);
```

# Example: Multiple outputs (Python)

```
Class Process(beam:DoFn):  
    def process(self, element):  
        try:  
            yield value.TaggedOutput('success', validateElement(element))  
        Except:  
            yield value.TaggedOutput('failure', element)  
success, failures = input | beam.ParDo(Process()).withOutputs("success", "failure")
```

# Example: Multiple outputs (Go)

```
success, failures := beam.ParDo2(s,  
    func(input: InputT, emitSuccess, emitFailure func(InputT)) {  
        if (validateElement(input) != nil) {  
            emitSuccess(input)  
        } else {  
            emitFailure(input)  
        }  
    }
```




# 03 Windowing & Time


*Where* in event time?



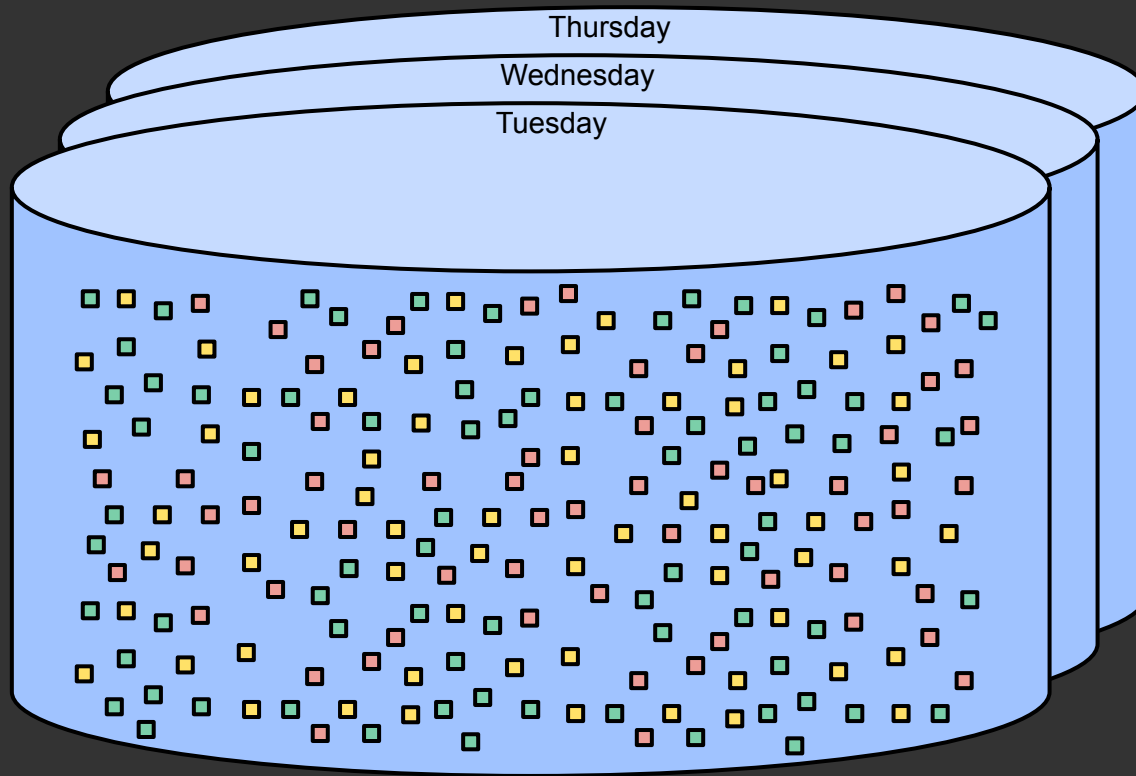
# Data...




...can be big...



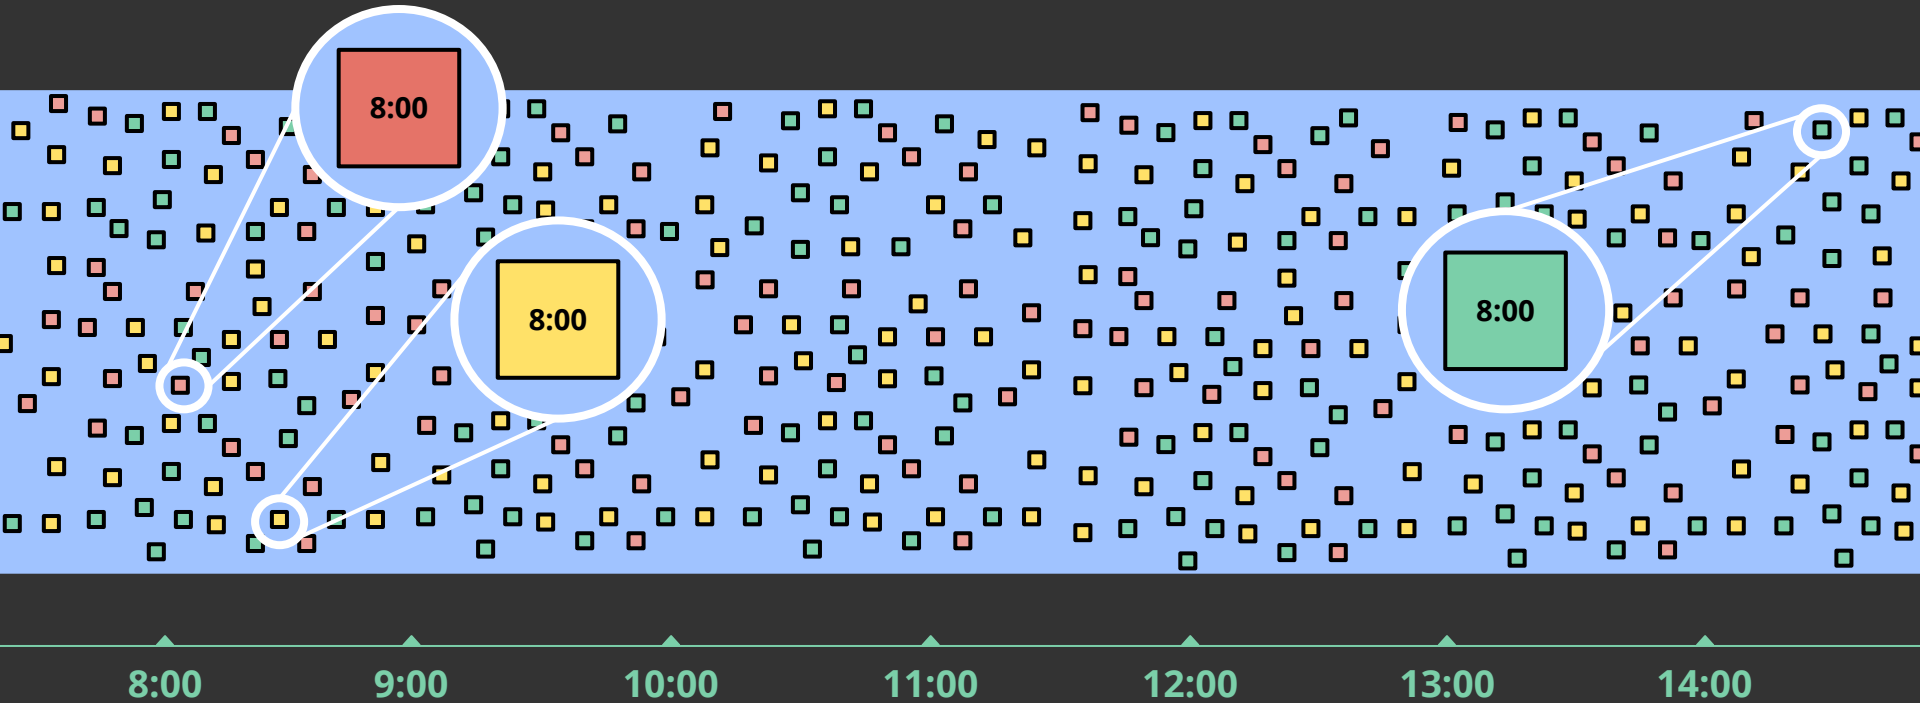
...really, really big...




...maybe infinitely big...




...with unknown delays.




# Element-wise transforms



# Grouping via processing-time windows




# Grouping via event-time windows





# What is windowing?


**Windowing** divides data into event-time-based finite chunks.



A windowing function computes which window(s) an element belongs to. Temporal functions can be parameterized with **duration** and **frequency**.

Often required when doing aggregations over unbounded data.

# What about data-dependent windowing?




Unique per key - you can't know a priori when a session ends, so the windowing function is now also parameterized by state.



# 04 Triggers & Streaming

*When* in processing time are results emitted?

# Streaming: Unbounded PCollections




Windowing specifies *where* events are aggregated in event time, but *when* are events emitted in processing time?

## *When*: triggering at the watermark

```
PCollection<KV<String, Integer>> scores = input
    .apply(Window.into(FixedWindows.of(Minutes(2))))
    .apply(Sum.integersPerKey());
```

# When: triggering at the watermark



Triggers control **when** the aggregation is output.


The default is “*when the watermark passes the end of the window*”.

This is the same as “*when we estimate the window is complete*”

## When: early & late firings

```
PCollection<KV<String, Integer>> scores = input
    .apply(Window.into(FixedWindows.of(Minutes(2))
        .triggering(AtWatermark()
            .withEarlyFirings(AtPeriod(Minutes(1)))
            .withLateFirings(AtCount(1))))))
    .apply(Sum.integersPerKey());
```

# When: early & late firings



Speculative triggers provide **early** updates *before* the watermark passes.

Watermark triggers provide **on-time** updates when input is believed complete.

Late triggers provide **late** updates when data arrive *after* the watermark (late data).



# Other kinds of triggers

## Element Count

Output after at least N elements

## Processing Time

Output after at least N minutes

## Combinators

Early/on-time/late

After all of these

After any of these

After each of these in order

etc.

Together these can be used for fine-grained control of output

For example:

- Early: every minute
- On-Time: when watermark predicts the window is complete
- Late: after every element



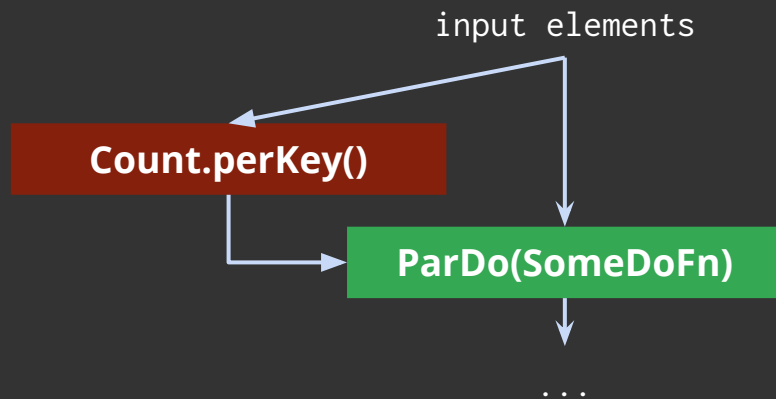
# 05 Side Inputs

# Side inputs

ParDos can receive extra inputs “on the side”

For example broadcast the count of elements to the processing of each element

Side inputs are computed (and accessed) per-window



# Example: ParDo with side inputs

```
PCollection<String> words = ...; // the input PCollection
PCollection<Integer> wordLengths = ...; // map words to their lengths

// Create a PCollectionView (singleton in this case).
// See also View.asList, View.asMap, etc.
final PCollectionView<Integer> maxWordLengthView =
    wordLengths.apply(Combine.globally(new Max.MaxIntFn()).asSingletonView());

// Apply a ParDo that takes maxWordLengthView as a side input, and left pads words.
PCollection<String> rightPaddedWords = words.apply(ParDo
    .withSideInputs(maxWordLengthView).of(new DoFn<String, String>() {
        @ProcessElement
        public void processElement(ProcessContext c) {
            int length = c.sideInput(maxWordLengthView);
            String format = "%1-" + length + "s";
            c.output(String.Format(format, c.element()));
        }
    }));
```

# Scala API! Scio

```
sc.textFile(input)
  .map { w => w.trim } // trim whitespace.
  .filter { w => w.nonEmpty } // filter out empty lines.
  .flatMap(_.split("[^a-zA-Z']+" ).filter(_.nonEmpty)) // split lines
  .countByValue
  .map(t => t._1 + ":" t._2) // format output word:count
  .saveAsTextFile(output)
```

Thank you!